



<CAPS>

Session 1 – Foundations Bootcamp

Leo Klenner, Henry Fung, Cory Combs

Outline

1. Admin
2. Introduction to CAPS
3. Case Study on Learning in Human Teams
4. Reinforcement Learning Basics



Admin

- > Two sessions per week. Always 6-7:30 pm
 - > Wednesday in Nitze 507
 - > Friday in BOB 736
- > Skills course policy
 - > Attend all session, get transcript certification
- > Course website at github.com/capsseminar



Introduction to CAPS

- > What skills do we need to guide the implementation of artificial intelligence technologies into the international arena?
 - > Explore a relevant skill set through the fictional **Forward Deployed Policy Engineer**



Training as a Forward Deployed Policy Engineer

- > Differentiate levels of autonomous decision making based on latent performance factors
- > Navigate core algorithmic learning architectures, such as reinforcement learning
- > Transform the specification of a learning algorithm into decision-relevant intelligence
- > Infer a system's underlying specifications from its observed behavior
- > Find misspecifications in a software architecture based on fuzzy end-user feedback
- > Identify performance-relevant links in complex human-machine ecosystems
- > Communicate technical topics to senior policy makers and policies to engineers



Case Study: Learning in a COIN Team

- > Army captains Jason Emory and Mark Nutsch
 - > Deployed as one the first Special Forces detachments to Afghanistan in 2001
 - > Our goal is to understand how they relate their training to their deployment
- > Understand general conceptual trade-offs that also apply to machine learning
- > Coverage of training, modes of learning, relevance of training data set



Markov Decision Processes

- > **Agent** senses its current **state** s_t , chooses an **action** a_t from a set of all possible actions A , and performs it.
- > In response, **environment** “gives” the agent two things:
 - > Immediate **reward** r_t
 - > Subsequent state s_{t+1}
- > The task of the agent is to learn a “**policy**” that dictates the choice of its action a_t based on its current observed state s_t :

$$\pi(s_t) = \forall s_t \in S$$



The Cumulative Discounted Reward

- > Overtime the agent gets a sequence of rewards for each action that it takes:

$$V^{\pi}(s_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

- > The agent starts in s_t , and repeatedly uses policy π to select its actions.
- > The agent gets a cumulative reward V from starting in s_t and following π thereafter.
- > What is the optimal policy the robot should learn?



The Learning Task

- > We require the agent to learn a policy π that maximizes its cumulative reward for all s . we call such a policy π^* the “optimal policy”.


$$\pi^* = \underset{\pi}{argmax} V^{\pi}(s)$$

- > In other words, the agent must follow a sequence of actions, dictated by policy π^* , that allows it to receive the maximum possible cumulative reward $V^*(s)$.



Components of the Six-Grid World

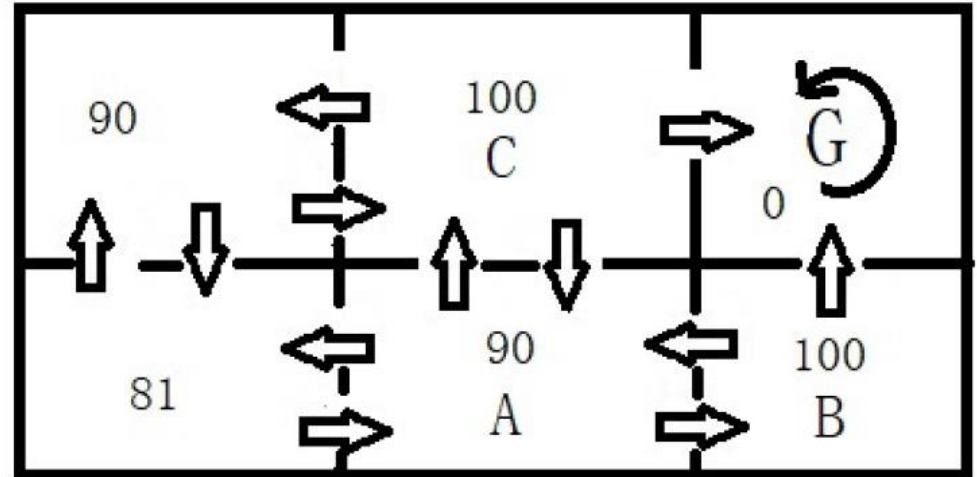

- > G is the goal state
- > The discount factor $\gamma = 0.9$
- > Arrows represent the possible actions that the agent can perform in each state
- > The numbers on the arrows are the immediate rewards
- > **The agent's learning task:** starting from an arbitrary state, move to G on the shortest path



Six-Grid World $V^*(s)$ values

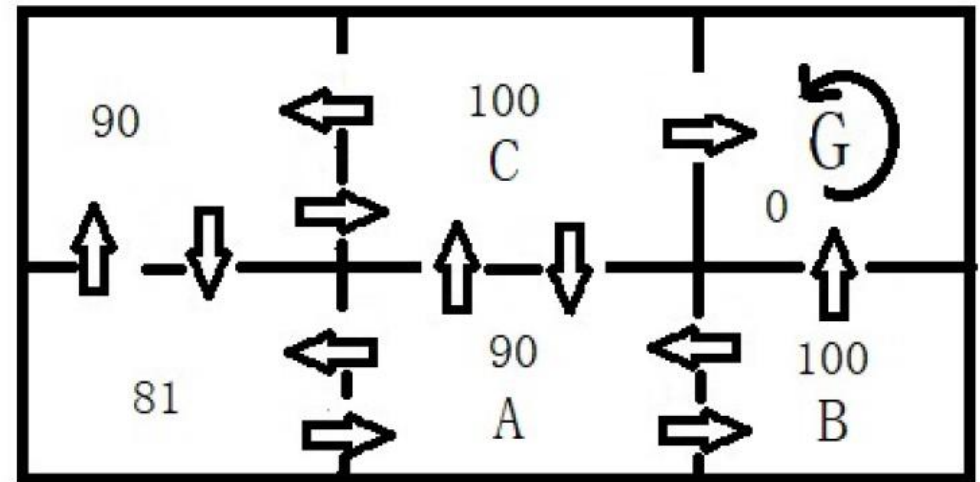
- > The numbers labelled in each grid in the bottom figure represents the largest possible cumulative reward that the agent receives if it follows the optimal policy
- > What is an optimal policy if the agent starts in A?

$$V^\pi(s = A) = 0 + 0.9 * 100 = 90$$



$V^*(s)$ as an Evaluation Function

- > It is difficult for the agent to learn π^* since it can only observe a sequence of immediate rewards
- > Trick: “Learn” $V^*(s)$ by playing the game many times. Then, use $V^*(s)$ to develop a preference over different states
- > Suppose the agent has learned $V^*(s)$, the values in the figure. If the agent is in A, what is the “best” action that it should perform?



$$\pi^*(s) = \operatorname{argmax}[r(s, a) + \gamma V^*(\delta(s, a))]$$

Problems with Using $V^*(s)$

- > What if we don't have previous knowledge of the immediate reward and the subsequent state for each (s, a) ?

$$\pi^*(s) = \operatorname{argmax}[r(s, a) + \gamma V^*(\delta(s, a))]$$


- > Solution: Need to do some clever rewriting and massage the above equation to bypass this problem.



The Q-function

- > Instead of evaluating each state, we now evaluate each state-action pair”. How good it is to perform action a in state s ?
- > The value Q is the immediate reward, plus the discounted cumulative reward that the robot gets by following the optimal policy thereafter.
- > Example:


$$Q(s = A, a = \text{right}) = 0 + 0.9 * 100 = 90$$



$$Q(s, a) = r(s, a) + \gamma V^*(\delta(s, a))$$

Selecting Optimal Actions with Q-values

- > To choose the optimal action a in state s , we simply select the action with the highest Q-value.
- > What is the agent's action in state A?
- > In state B?
- > The agent doesn't need prior knowledge of the reward and the subsequent state that it will receive from the environment by performing an action to select its optimal action a^* , once it learns Q .



$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$

The Bellman Equation

- > The Q-function can be rewritten as the Bellmann Equation:

$$Q(s, a) = r(s, a) + \gamma \max_{a'} Q(\delta(s, a), a')$$

- > Learning Q corresponds to learning the optimal policy.
- > Armed with the Bellman Equation, we can use a *Value Iteration* algorithm to estimate the Q values (Session3).
- > What is $Q(s=A, a=\text{right})$ from the Bellmann Equation?

