



<CAPS>

Session 3 – From Learning Architectures to Decision Profiles

Leo Klenner, Henry Fung, Cory Combs

From Deduction to Induction

Deduction

Learning architectures \rightarrow decision profiles

Learning architectures $\overset{?}{\leftarrow}$ decision profiles

Induction



Review of Sessions 1 and 2

- > Architectures for decision making
 - > **Learning in human teams** (supervision, interaction)
 - > **Reinforcement learning** (online-offline, learning with state-action pairs)
 - > **Rule-based systems** (Finite State Machines, conditional logic)
 - > **Multi-agent learning** (joint-separate, “Policy Geometry”)

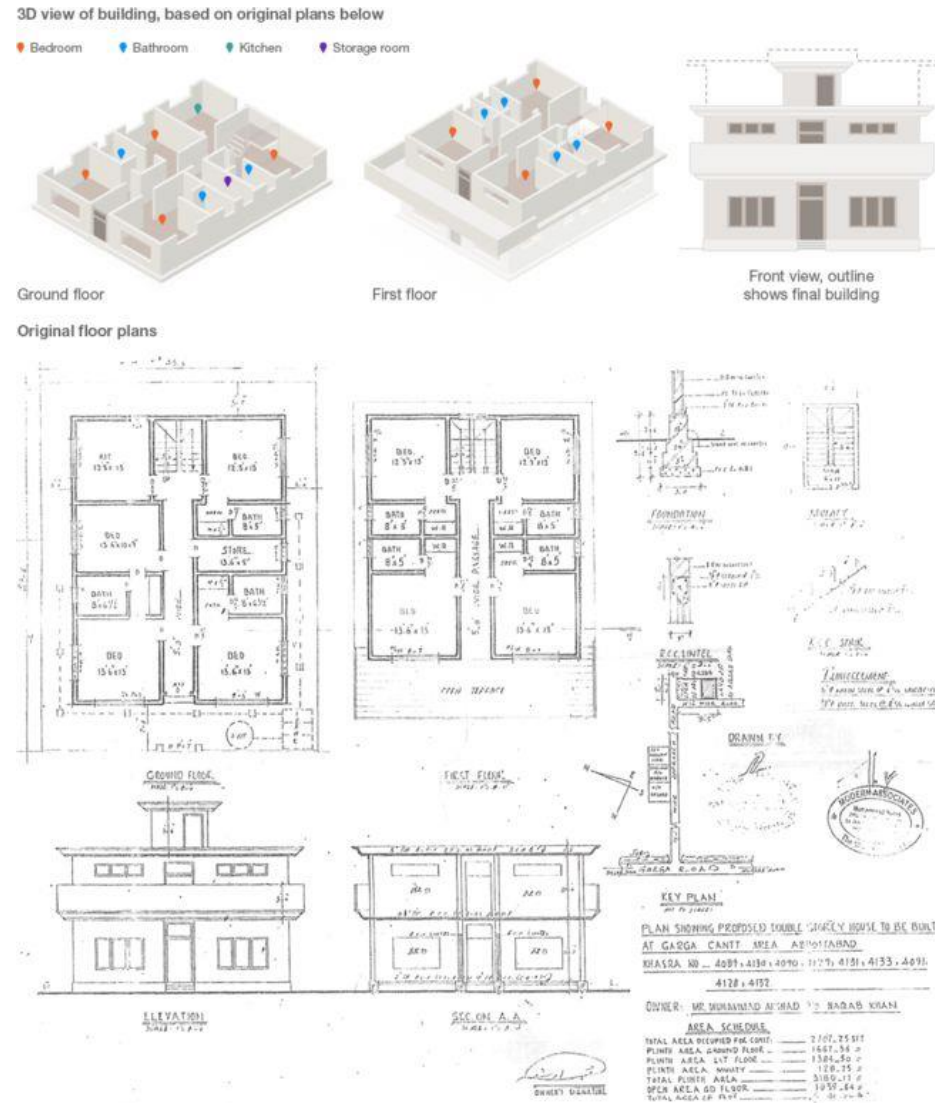


Outline

1. Exercise on Rule-Based Systems
2. Recap on Reinforcement Learning
3. Q-Learning
4. Case on Inferring Learning Architectures from Observed Behavior




Location X



Review of Reinforcement Learning Basics

- > RL is suitable to a class of problems known as the Markov Decision Process (MDP)
 - > Outcome of applying action to any state depends only on this action a and state s (and not on preceding a or s)
 - > Information available to the agent: rewards and subsequent state for each state-action transition




> *What are the components of a MDP?*



What is the Total (Cumulative) Reward?

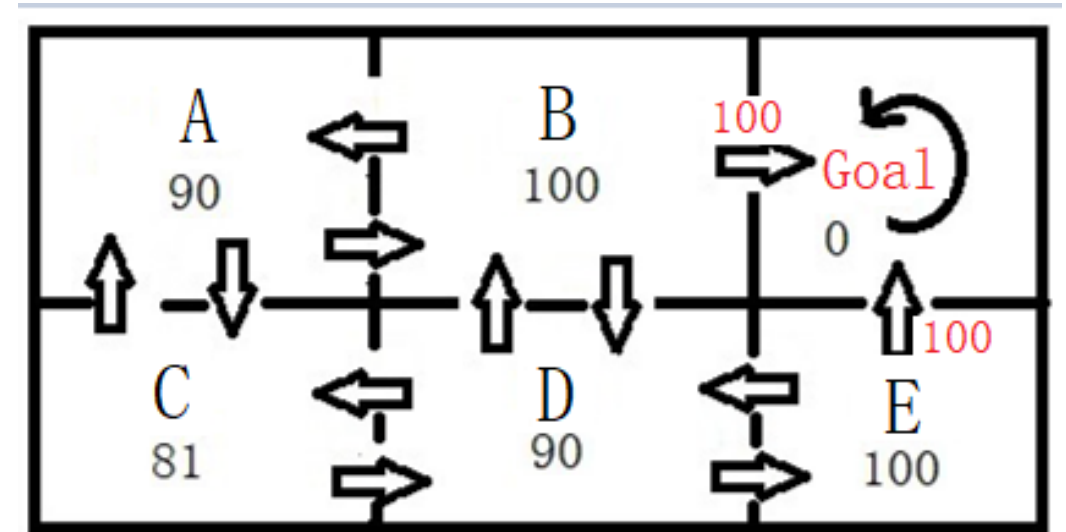
- > After performing action a in state s (state action pair s, a):
 - > The agent receives immediate reward of either 0 or 100
 - > The agent goes to the next state
- > The action that the agent chooses in every state s is dictated by a policy π
- > $\pi(s_t) = a_t$
- > The total reward that the agent gets (from π) is the sum of the sequence of rewards that it gets by following π , discounted by γ



$$V^\pi(s_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

What is the Agent's Goal?

- > What is the goal of an agent that uses RL?
 - > Pick the best policy that will maximize the total rewards that the agent receives
- > How does it do that?
 - > The agent needs to evaluate the “goodness” of each state. How good is it to be in Grid A, in Grid E?
 - > This evaluation can be done if the agent knows the maximum total reward that it can achieve in each state
 - > We call this the “optimal value function” $V^*(s)$
 - > If we have good estimates of $V^*(s)$, then we can choose the optimal policy




$$V^*(s = A)$$

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} [r(s, a) + \gamma V^*(\delta(s, a))]$$



The Q-Function

- > Rather than evaluating the “goodness” of each state, we can also evaluate the goodness of each state-action pair
 - > How good it is for me to perform RIGHT in state D (as oppose to RIGHT) in state B?
 - > Compare the Q value of doing RIGHT in D and RIGHT in B
 - > Q value is the maximum achievable total reward that the robot get by performing a in s , and then following optimal policy thereafter



$$Q(s, a) = r(s, a) + \gamma V^*(\delta(s, a))$$

$$Q(s = D, a = right) = 0 + 0.9 * 100 = 90$$




The Bellman Equation

- > The Q-function can be rewritten as the Bellmann Equation:

$$Q(s, a) = r(s, a) + \gamma \max_{a'} Q(\delta(s, a), a')$$

- > Learning Q corresponds to learning the optimal policy.
- > What is $Q(s=D, a=RIGHT)$ from the Bellmann Equation?
- > Learning Q corresponds to learning the optimal policy. Assuming Q-values are known, what is the optimal strategy starting from A?



$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$



Learn $V^*(s)$ or $Q(s, a)$?


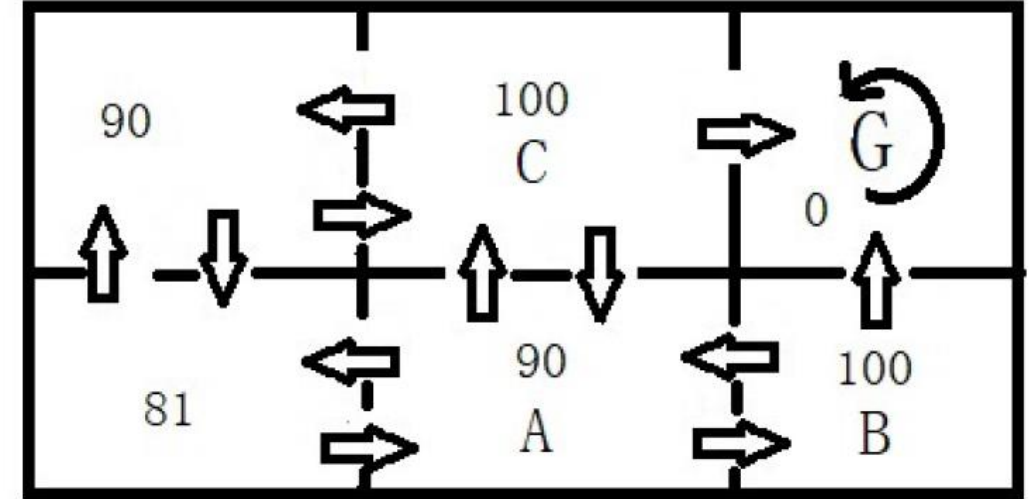
- > **Approach 1:** If we know the true $V^*(s)$ values, we know the optimal policy:

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} [r(s, a) + \gamma V^*(\delta(s, a))]$$

- > **Approach 2:** If we know the true Q-values, we know the optimal policy:

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$

- > Which approach should we choose?



Model-based vs. Model-free Learning

Model-Based Learning

- > Model of the environment : reward and state-transition need to be known
- > The agent interacts with environment, and from the “history” of its interactions, the agent will approximate the reward and state transition
- > After that, the agent learns the $V^*(s)$ using an algorithm called “Value iteration”. Once the $V^*(s)$ values are known, the agent can find the optimal policy

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} [r(s, a) + \gamma V^*(\delta(s, a))]$$

Model Free Learning

- > The agent will not try to “guess” the reward and state transition functions. It will only “experience” them as it explores the environment
- > By trail and error, the robot discovers what are the good and bad actions in each state
- > Uses “Q learning” to approximate the $Q(s, a)$ values and find optimal policy

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$



The Q-Learning Algorithm

> Assumptions:

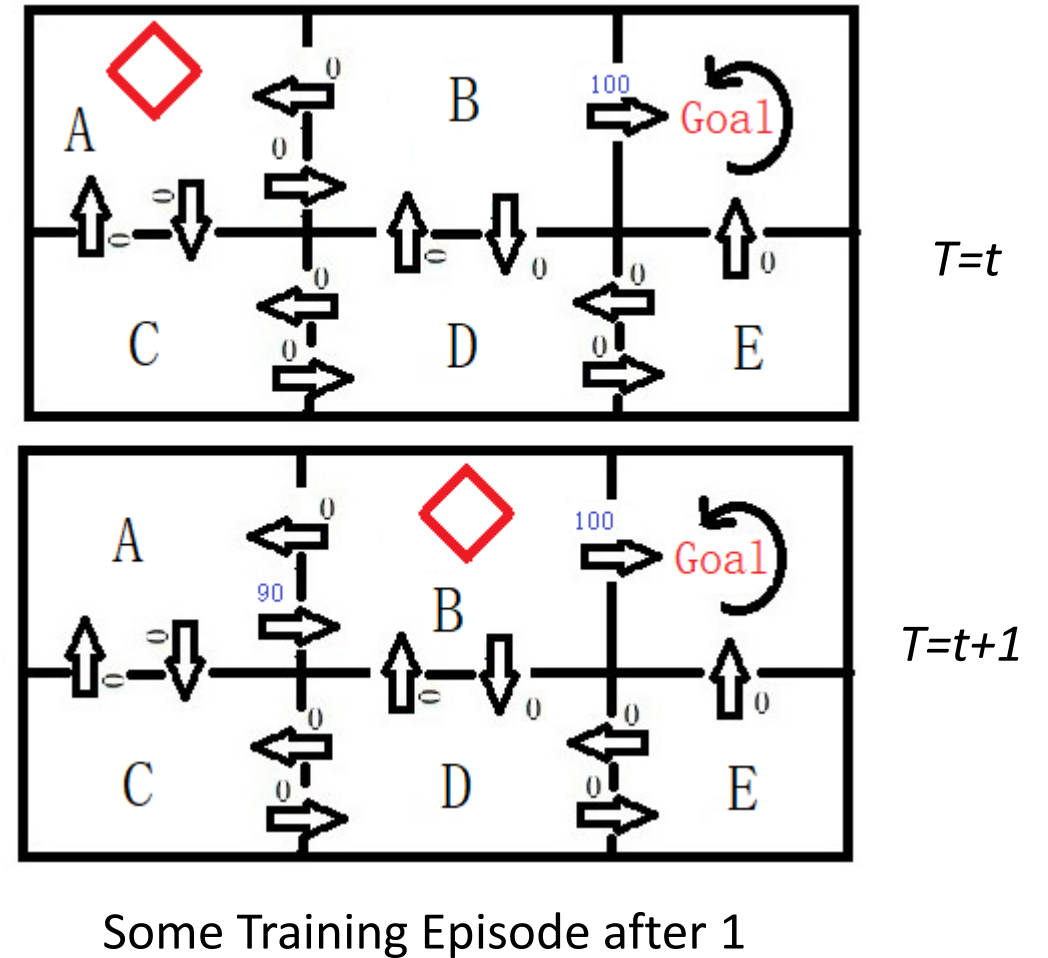
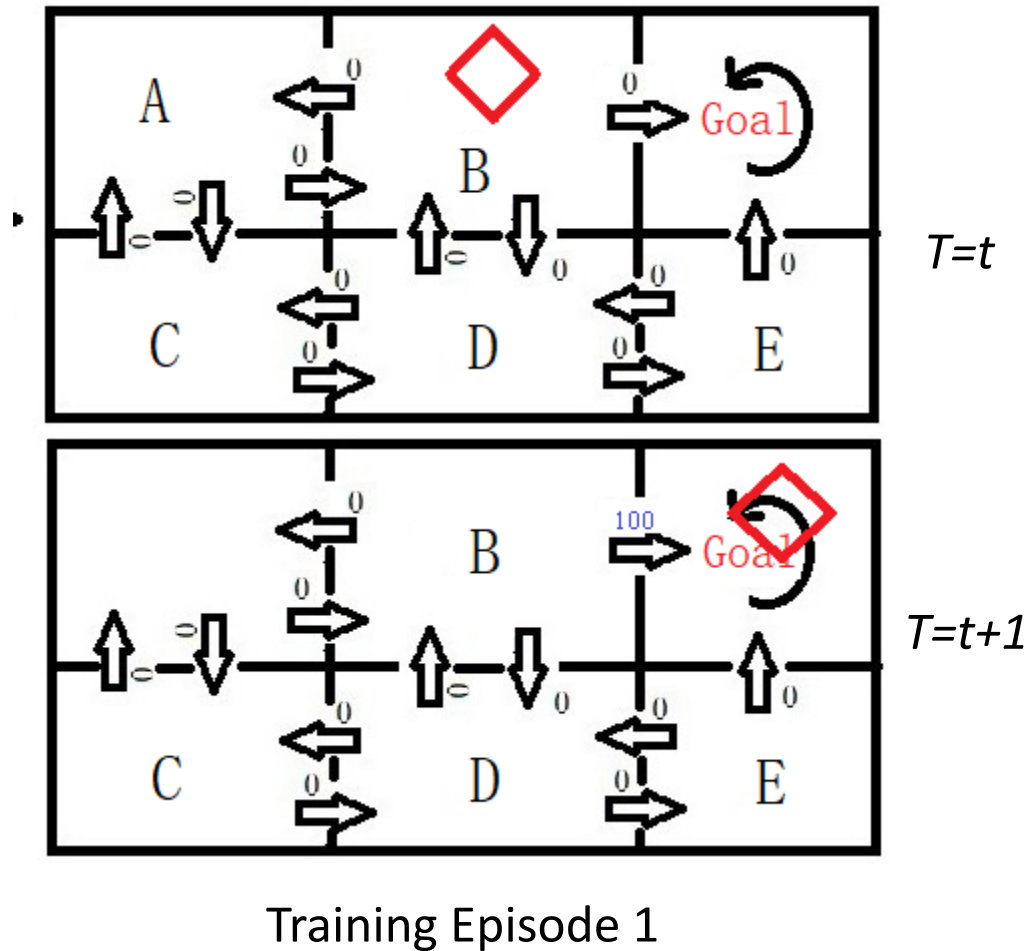
- > (A1) Reward and state-transitions functions are unknown but deterministic.
- > (A2) Robot visit every possible state with non-zero frequency overtime.
- > (A3) Rewards are bounded (not infinite).

> Algorithm:

1. Initially fill the table of \hat{Q} (estimated Q-values) with zeros
2. Observe current state s
3. Select an action a and execute it
4. Receive immediate reward r
5. Observe the new state s'
6. Update the table entry for $\hat{Q}(s, a)$ using the Bellman equation
7. Set state s' to s , repeat 2-6 until \hat{Q} converges (new \hat{Q} values do not change much from old values)



Example: Information Propagation



Strategies for Experimentation

- > The Q-learning algorithm does not specify how the agent chooses its actions
- > If the agent *always* selects the state-action pair with the highest estimated Q-values, then it will “overcommit” to its actions found in the early training sessions, while failing to explore other state-action pairs that might have even higher Q-values
- > Solution to Exploration vs. Exploitation dilemma:
 - > During early training episodes: All actions have a chance of being executed
 - > During the latter training episodes: Actions with higher \hat{Q} will be assigned with increasingly higher probabilities, so overtime, the agent will do more exploitation and less exploration




Decision Profiles in Unbounded Environments

Learning architectures $\overset{?}{\leftarrow}$ decision profiles

- > Does each architecture correspond to a distinctive decision profile?
- > How can we discern this decision profile?
- > What are the consequences, if we answer the first or second question negatively?



AV Flight Trajectory



- > Goerzen, et al. 2009. A Survey of Motion Planning Algorithms from the Perspective of Autonomous UAV Guidance



Toy Model of Learning Architectures

- > Does each architecture correspond to a distinctive decision profile?

Rule-based → precise solution, no adaptation


Reinforcement learner → multiple degrees of adaptation

Human → full adaptation and rule generation

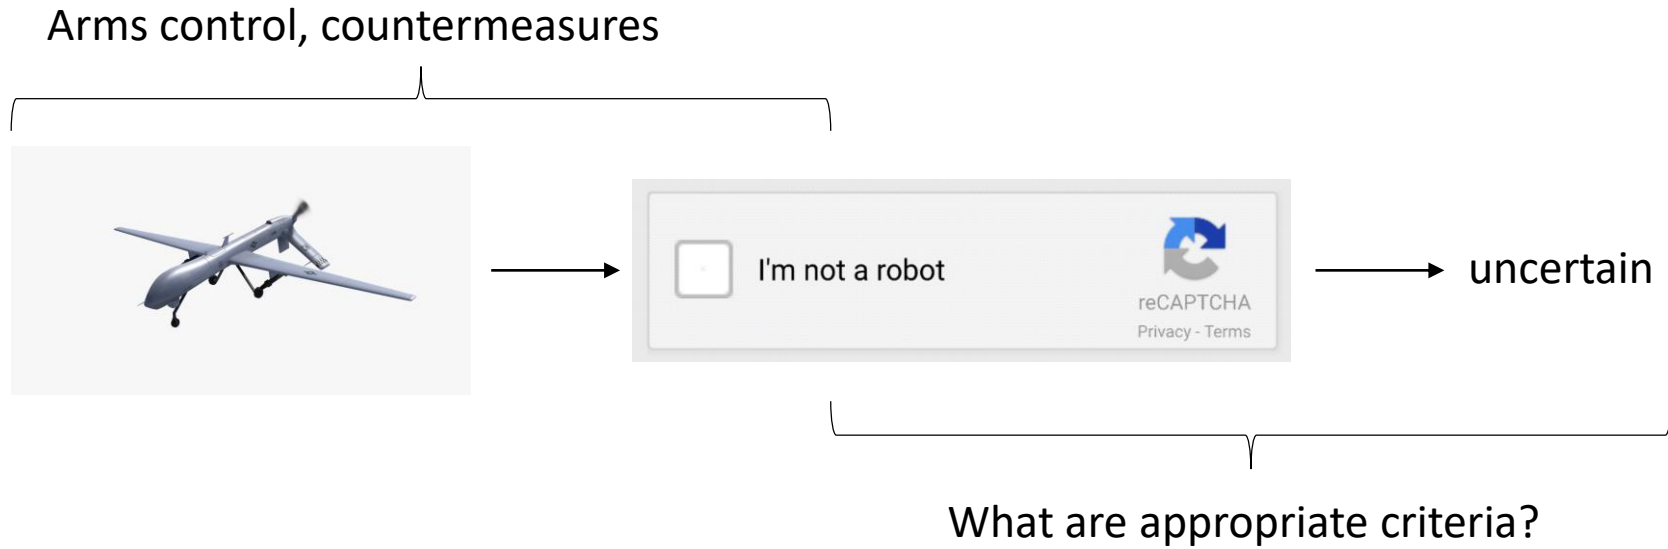
- > Assumes we have perfect information about the environment
- > Real-world is often too complex to justify such an assumption



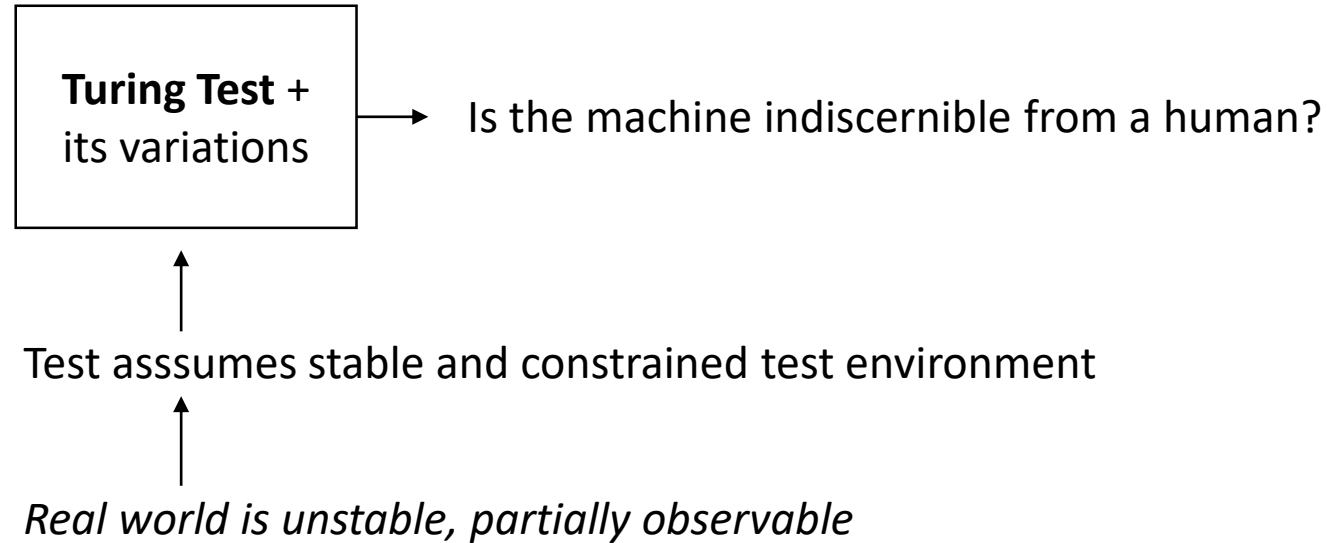
Possible Architectures for Trajectory



Why Care about Verification?



Differentiating Human from Machine



- > **Hypothesis:** A system that would not pass the Turing Test in theory might pass it in practice, because the test environment is fuzzy




Other Mechanisms of Verification

Performance	Worse than human	Human-level	Better than human
Verification	Easy ↓ Clear failure modes	Hard ↓ Internal factors: Unexpected decision patterns Unexpected failure modes External factors: Change in force structure Change in strategy Introduce disturbances in env.	Easy; depends ↓ Unparalleled performance



Petrov Verification




- > September 26, 1983
- > Five U.S. ICBMs surface on Soviet radar, station monitored by Stanislav Petrov
- > Petrov reasons that this is a false alarm, decides not to launch counter-attack



= Petrov Verification: combination of internal and external factors; educated guess *across* domains on limited information



Appendix on Countermeasures



- > Krishna et al. 2005. Parametric Control of Multiple Unmanned Air Vehicles over an Unknown Hostile Territory