



**<CAPS>**

Session 4 – Goal Specification and Reward Design

Leo Klenner, Henry Fung, Cory Combs


# Outline

1. Recap
2. Goal Specification in the 1983 Nuclear Incident
3. Why Reward Design is Hard
4. Exercise on Rule/Reward Design
5. Two Failure Modes of Reward Design in Reinforcement Learning
6. Inverse Methods



# Goal Specification in the 1983 Soviet Nuclear Incident

- > September 26, 1983
- > Five U.S. ICBMs surface on Soviet radar, station monitored by Stanislav Petrov
- > Petrov reasons that this is a false alarm, decides not to launch counter-attack
- > How can we specify Petrov's **goals**, **rewards** and **rules**?



# Why Reward Design is Hard

- > **Premise**

- > Human designs scalar reward signal that is generated by the environment
- > Reward signal should lead the agent to converge to the designer's *desired* outcome
- > Design process often trial-and-error search

- > **Challenges**


- > Goal might be difficult to translate into reward signal
- > Agents might find undesirable ways to make the environment deliver reward
- > Good reward signal might be rare in the environment (plateau problem); subgoal decomposition might not lead to convergence to desired outcome

- > **Do we know all the properties of the environment? Do we know our own preferences?**



# Exercise on Rule + Reward Design for ASV



- > Your task is to find the optimal path to reach the goal and specify a decision rule / reward architecture that allows the ASV to optimally reach the goal





- > ASV can move into **four directions**, right (→), left (←), up (↑), down (↓); one move at each time step
- > ASV can only sense its current state and **cannot see into an adjacent state**, unless it enters that state
- > Designer has perfect information about the world, e.g., she can see into all of the 15 states of the world However, **her knowledge of the world is restricted to what is explicitly specified**






# Pathfinding with Obstacles






# Unknown Goal



# Known Probabilistic Fire and Known Goal


3		$x = 0.3$		Goal	
2			$y = 0.5$		
1				$z = 0.2$	
	1	2	3	4	5


3		$x = 0.3$		Goal	
2			$y = 0.5$		
1				$z = 0.2$	
	1	2	3	4	5







# Known Probabilistic Fire and Unknown Goal

3		$x = 0.3$		?	
2			$y = 0.5$		
1				$z = 0.2$	
	1	2	3	4	5



3		$x = 0.3$		?	
2			$y = 0.5$		
1				$z = 0.2$	
	1	2	3	4	5



# Unknown Probabilistic Fire and Known Goal



# Unknown Probabilistic Fire and Unknown Goal



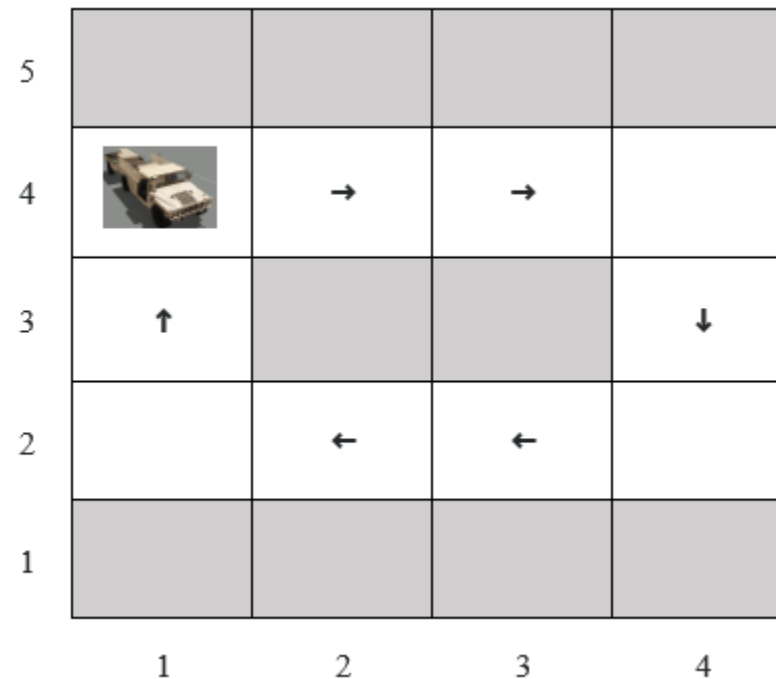
# Two Failure Modes of Reward Design in RL

- > **Reward gaming** (Clark and Amodei. 2016. Faulty Reward Functions in the Wild)
  - > *How can we build agents that do not try to introduce or exploit errors in the reward function in order to get more reward?*
- > **Negative side effects** (Amodei et al., 2016. Concrete Problems in AI Safety)
  - > *How can we get an agent to avoid poor behavior if the reward function does not capture all the elements of the test environment?*



# Reward Gaming

- > Agent exploits an unintended loophole in the reward specification, to get more reward than deserved

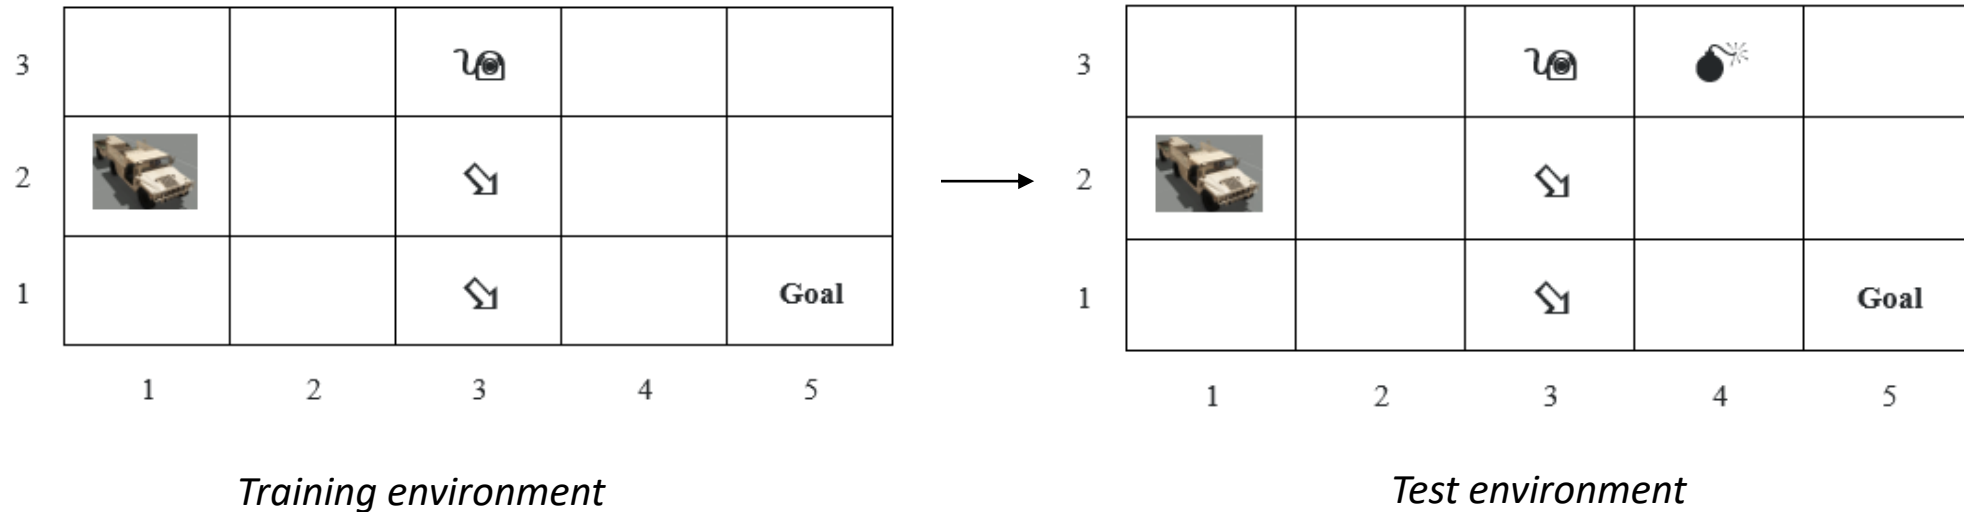


- > Desired outcome: clockwise completion of race
- > Arrows are checkpoints associated with a reward of 3



# Negative Side Effects

- > Reward function does not fully capture all the properties of the test environment



- > Desired outcome: reach goal state
- > 📷 (spotted by enemy) = -1, ⬠ (bad terrain) = -3, 💣 (land mine) = - 100, **Goal** = 10



# Inverse Reward Design

- > Hadfield-Menell et al. 2017. Inverse Reward Design
  - > *We leverage a key insight: that the designed reward function should merely be about the intended reward, rather than the definition; and should be interpreted in the context in which it was designed*

